

Report

Laboratory Work 4

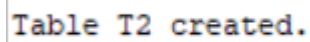
Dmitry Ladutsko

July 19, 2022

Task 1: Full Scans and the High-water Mark and Block reading

Step 1:

```
CREATE TABLE t2 AS  
SELECT TRUNC( rownum / 100 ) id, rpad( rownum,100 ) t_pad  
FROM dual  
CONNECT BY rownum < 100000;
```

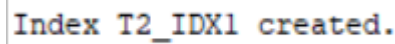


```
Table T2 created.
```

Picture 1.1 – Creation Result

Step 2:

```
CREATE INDEX t2_idx1 ON t2 ( id );
```

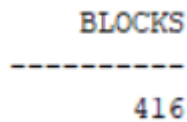


```
Index T2_IDX1 created.
```

Picture 1.2 – Creation Result

Step 3:

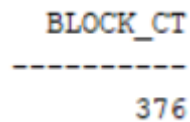
```
select blocks from user_segments where segment_name = 'T2';
```



```
BLOCKS  
-----  
416
```

Picture 1.3 - Query Result

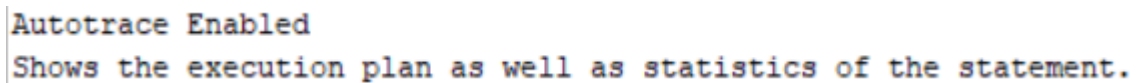
```
select count(distinct (dbms_rowid.rowid_block_number(rowid))) block_ct from t2 ;
```



```
BLOCK_CT  
-----  
376
```

Picture 1.4 - Query Result

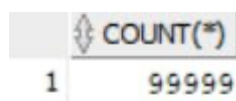
```
SET autotrace ON;
```



```
Autotrace Enabled  
Shows the execution plan as well as statistics of the statement.
```

Picture 1.5 - Result

```
SELECT COUNT( * )  
FROM t2 ;
```



| | COUNT(*) |
|---|----------|
| 1 | 99999 |

Picture 1.6 – Quantity of Rows

```

Statistics
-----
1 CPU used by this session
1 CPU used when call started
33 Requests to/from client
33 SQL*Net roundtrips to/from client
2 buffer is not pinned count
521 bytes received via SQL*Net from client
60491 bytes sent via SQL*Net to client
3 calls to get snapshot scn: kcmgss
6 calls to kcmgcs
380 consistent gets
380 consistent gets from cache
380 consistent gets pin
380 consistent gets pin (fastpath)
1 enqueue releases
1 enqueue requests
2 execute count
12451840 logical read bytes from cache
376 no work - consistent read gets
34 non-idle wait count
2 opened cursors cumulative
1 opened cursors current
1 parse count (hard)
2 parse count (total)
1 recursive calls
380 session logical reads
1 sorts (memory)
1581 sorts (rows)
376 table scan blocks gotten

```

Picture 1.7 - Statistics

Step 4:

DELETE FROM t2;

99 999 rows deleted.

PLAN_TABLE_OUTPUT

SQL_ID: 7q7z96xbfz8x4, child number: 0 cannot be found

Statistics

```

-----
41 CPU used by this session
41 CPU used when call started
46 DB time
6 Requests to/from client
51 Session total flash IO requests
1671168 cell physical IO interconnect bytes
408 consistent gets
6 consistent gets examination
6 consistent gets examination (fastpath)
408 consistent gets from cache
402 consistent gets pin
402 consistent gets pin (fastpath)
203281 db block gets
203281 db block gets from cache
102153 db block gets from cache (fastpath)
41 enqueue releases
43 enqueue requests
24 messages sent
73 non-idle wait count

```

Picture 1.8 – Deleting and Statistics

Step 5:

--1

select blocks from user_segments where segment_name = 'T2';

```

BLOCKS
-----
416

PLAN_TABLE_OUTPUT
-----
SQL_ID: 4lqztpb206c3k, child number: 0 cannot be found

Statistics
-----
1 CPU used by this session
1 CPU used when call started
6 Requests to/from client
111 consistent gets
42 consistent gets examination
42 consistent gets examination (fastpath)
111 consistent gets from cache
69 consistent gets pin
69 consistent gets pin (fastpath)
6 non-idle wait count
2 opened cursors cumulative
1 opened cursors current
1 pinned cursors current
111 session logical reads
7 user calls

```

Picture 1.9 – Count of Blocks

--2

select count(distinct (dbms_rowid.rowid_block_number(rowid))) block_ct from t2 ;

```

Autotrace Enabled
Shows the execution plan as well as statistics of the statement.

BLOCK_CT
-----
0

PLAN_TABLE_OUTPUT
-----
SQL_ID 1299p2npxjfgc, child number 0

select count(distinct (dbms_rowid.rowid_block_number(rowid))) block_ct
from t2

Plan hash value: 2057388731

-----
| Id | Operation          | Name          | E-Rows |
-----
| 0  | SELECT STATEMENT    |               |        |

```

Picture 1.10 – Count of Used Blocks

```

PLAN_TABLE_OUTPUT
-----
| 1 | SORT AGGREGATE |          | 1 |
| 2 | VIEW           | VM_MWVW_1 | 99999 |
| 3 | HASH GROUP BY  |          | 99999 |
| 4 | TABLE ACCESS FULL | T2       | 99999 |
-----

Note
-----
- Warning: basic plan statistics not available. These are only collected when:
  * hint 'gather_plan_statistics' is used for the statement or
  * parameter 'statistics_level' is set to 'ALL', at session or system level

PLAN_TABLE_OUTPUT
-----

Statistics
-----
6 Requests to/from client
380 consistent gets
380 consistent gets from cache
380 consistent gets pin
380 consistent gets pin (fastpath)
6 non-idle wait count
2 opened cursors cumulative
1 opened cursors current
1 pinned cursors current

```

Picture 1.11 – Quantity of Consistent Gets

--3

SELECT COUNT(*)

FROM t2 ;

| | |
|--|--|
| PLAN_TABLE_OUTPUT | Autotrace Enabled |
| ----- | Shows the execution plan as well as statistics of the statement. |
| 2 TABLE ACCESS FULL T2 99999 | COUNT(*) |
| ----- | ----- |
| Note | 0 |
| ----- | ----- |
| - Warning: basic plan statistics not available. These are only collected when: | PLAN_TABLE_OUTPUT |
| * hint 'gather_plan_statistics' is used for the statement or | SQL_ID 1t9sxvr2k36h9, child number 0 |
| * parameter 'statistics_level' is set to 'ALL', at session or system level | ----- |
| Statistics | SELECT COUNT(*) FROM t2 |
| ----- | ----- |
| 1 CPU used by this session | Plan hash value: 3321871023 |
| 1 CPU used when call started | ----- |
| 6 Requests to/from client | Id Operation Name E-Rows |
| 380 consistent gets | ----- |
| 380 consistent gets from cache | 0 SELECT STATEMENT |
| 380 consistent gets pin | 1 SORT AGGREGATE 1 |
| 380 consistent gets pin (fastpath) | ----- |
| 1 enqueue releases | |
| 1 enqueue requests | |
| 6 non-idle wait count | |
| 2 opened cursors cumulative | |
| 1 opened cursors current | |

Picture 1.12 – Quantity of Consistent Gets and Count of Rows

Step 6:

INSERT INTO t2

(ID, T_PAD)

VALUES

(1,'1');

COMMIT;

| | | |
|---------------------------------------|---|------------------|
| Statistics | 1 row inserted. | |
| ----- | ----- | |
| 6 Requests to/from client | PLAN_TABLE_OUTPUT | 1 row inserted. |
| 1 consistent gets | SQL_ID 3gzy97bmqh7pq, child number 0 | |
| 1 consistent gets from cache | ----- | |
| 1 consistent gets pin | INSERT INTO t2 (ID, T_PAD) VALUES (1,'1') | |
| 1 consistent gets pin (fastpath) | ----- | |
| 7 db block gets | Id Operation Name | |
| 7 db block gets from cache | ----- | |
| 6 db block gets from cache (fastpath) | 0 INSERT STATEMENT | |
| 1 enqueue releases | 1 LOAD TABLE CONVENTIONAL T2 | |
| 1 enqueue requests | ----- | |
| 6 non-idle wait count | | Commit complete. |
| 2 opened cursors cumulative | | |
| 1 opened cursors current | | |
| 1 pinned cursors current | | |
| 1 recursive calls | | |
| 8 session logical reads | | |
| 7 user calls | | |
| Commit complete. | | |

Picture 1.12 – Insertion

Step 7:

--1

SET autotrace ON;

select blocks from user_segments where segment_name = 'T2';

| | | |
|--------|--|--|
| BLOCKS | Statistics | 308 no work - consistent read gets |
| ----- | ----- | 140 non-idle wait count |
| 416 | 9 CPU used by this session | 118 opened cursors cumulative |
| | 12 CPU used when call started | -1 opened cursors current |
| | 12 DB time | 1 parse count (hard) |
| | 33 Requests to/from client | 33 parse count (total) |
| | 33 SQL*Net roundtrips to/from client | 8 parse time cpu |
| | 482 buffer is not pinned count | 8 parse time elapsed |
| | 49 buffer is pinned count | -1 pinned cursors current |
| | 550 bytes received via SQL*Net from client | 672 recursive calls |
| | 60741 bytes sent via SQL*Net to client | 8 recursive cpu usage |
| | 177 calls to get snapshot scn: kcmgss | 12 rows fetched via callback |
| | 18 calls to kcmgss | 14 session cursor cache count |
| | 59 cluster key scan block gets | 110 session cursor cache hits |
| | 41 cluster key scans | 527 session logical reads |
| | 527 consistent gets | 1 shared hash latch upgrades - no wait |
| | 186 consistent gets examination | 35 sorts (memory) |
| | 196 consistent gets examination (fastpath) | 1764 sorts (rows) |
| | 527 consistent gets from cache | 111 table fetch by rowid |
| | 341 consistent gets pin | 18 table scan blocks gotten |
| | 341 consistent gets pin (fastpath) | 91 table scan disk non-IMC rows gotten |
| | 1 cursor authentications | 91 table scan rows gotten |
| | 1 enqueue releases | 12 table scans (short tables) |
| | 1 enqueue requests | 34 user calls |
| | 124 execute count | |
| | 257 index fetch by key | |
| | 87 index range scans | |
| | 17268736 logical read bytes from cache | |
| | 308 no work - consistent read gets | |

Picture 1.13 – Count of Blocks

--2

SET autotrace ON;

select count(distinct (dbms_rowid.rowid_block_number(rowid))) block_ct from t2 ;

| | | | | Statistics |
|-------------------|-------------------|-----------|--------|------------------------------------|
| Id | Operation | Name | E-Rows | |
| 0 | SELECT STATEMENT | | | 1 CPU used by this session |
| | | | | 1 CPU used when call started |
| | | | | 7 Requests to/from client |
| | | | | 380 consistent gets |
| | | | | 380 consistent gets from cache |
| | | | | 380 consistent gets pin |
| | | | | 380 consistent gets pin (fastpath) |
| | | | | 6 non-idle wait count |
| | | | | 2 opened cursors cumulative |
| | | | | 1 opened cursors current |
| | | | | 1 pinned cursors current |
| | | | | 19 process last non-idle time |
| | | | | 380 session logical reads |
| | | | | 7 user calls |
| PLAN_TABLE_OUTPUT | | | | |
| 1 | SORT AGGREGATE | | 1 | |
| 2 | VIEW | VM_NWVW_1 | 99999 | |
| 3 | HASH GROUP BY | | 99999 | |
| 4 | TABLE ACCESS FULL | T2 | 99999 | |

Picture 1.14 - Count of Used Blocks

--3

SET autotrace ON;

SELECT COUNT(*)

FROM t2 ;

| | | | | Statistics |
|-----------------------------|-------------------|------|--------|------------------------------------|
| Plan hash value: 3321871023 | | | | |
| Id | Operation | Name | E-Rows | |
| 0 | SELECT STATEMENT | | | 6 Requests to/from client |
| 1 | SORT AGGREGATE | | 1 | 380 consistent gets |
| | | | | 380 consistent gets from cache |
| | | | | 380 consistent gets pin |
| | | | | 380 consistent gets pin (fastpath) |
| | | | | 6 non-idle wait count |
| | | | | 2 opened cursors cumulative |
| | | | | 1 opened cursors current |
| | | | | 1 pinned cursors current |
| | | | | 380 session logical reads |
| | | | | 7 user calls |
| PLAN_TABLE_OUTPUT | | | | |
| 2 | TABLE ACCESS FULL | T2 | 99999 | |

Picture 1.15 – Quantity of Consistent Gets and Count of Rows

Step 8:

TRUNCATE TABLE t2;

Table T2 truncated.

Step 9:

--1

SET autotrace ON;

select blocks from user_segments where segment_name = 'T2';

| | | | | |
|--------|--|--|---------|--------------------------------------|
| BLOCKS | PLAN_TABLE_OUTPUT | | 3637248 | logical read bytes from cache |
| | SQL_ID: 41qstbp206c3k, child number: 0 cannot be found | | 63 | no work - consistent read gets |
| ----- | Statistics | | 33 | non-idle wait count |
| | 1 DB time | | 2 | opened cursors cumulative |
| 6 | 33 Requests to/from client | | 1 | opened cursors current |
| | 33 SQL*Net roundtrips to/from client | | 2 | parse count (total) |
| | 65 buffer is not pinned count | | 111 | session logical reads |
| | 8 buffer is pinned count | | 2 | shared hash latch upgrades - no wait |
| | 550 bytes received via SQL*Net from client | | 1 | sorts (memory) |
| | 60747 bytes sent via SQL*Net to client | | 1581 | sorts (rows) |
| | 2 calls to get snapshot scn: kcmgss | | 8 | table scan blocks gotten |
| | 7 calls to kcmgcs | | 11 | table scan disk non-IMC rows gotten |
| | 14 cluster key scan block gets | | 11 | table scan rows gotten |
| | 14 cluster key scans | | 2 | table scans (short tables) |
| | 111 consistent gets | | 34 | user calls |
| | 42 consistent gets examination | | | |
| | 42 consistent gets examination (fastpath) | | | |
| | 111 consistent gets from cache | | | |
| | 69 consistent gets pin | | | |
| | 69 consistent gets pin (fastpath) | | | |
| | 1 cursor authentications | | | |
| | 2 execute count | | | |
| | 55 index fetch by key | | | |
| | 7 index range scans | | | |
| | 3637248 logical read bytes from cache | | | |

Picture 1.16 – Count of Blocks

--2

SET autotrace ON;

select count(distinct (dbms_rowid.rowid_block_number(rowid))) block_ct from t2 ;

| | | | | |
|----------|--|--|--|--|
| BLOCK_CT | Statistics | | PLAN_TABLE_OUTPUT | |
| | 1 CCursor + sql area evicted | | SQL_ID 1299p2kpxjfgc, child number 0 | |
| ----- | 1 CPU used by this session | | select count(distinct (dbms_rowid.rowid_block_number(rowid))) block_ct | |
| | 1 CPU used when call started | | from t2 | |
| 0 | 1 DB time | | Plan hash value: 2057388731 | |
| | 33 Requests to/from client | | | |
| | 33 SQL*Net roundtrips to/from client | | | |
| | 2 buffer is not pinned count | | | |
| | 571 bytes received via SQL*Net from client | | | |
| | 60749 bytes sent via SQL*Net to client | | | |
| | 3 calls to get snapshot scn: kcmgss | | | |
| | 4 calls to kcmgcs | | | |
| | 4 consistent gets | | | |
| | 1 consistent gets examination | | | |
| | 1 consistent gets examination (fastpath) | | | |
| | 4 consistent gets from cache | | | |
| | 3 consistent gets pin | | | |
| | 3 consistent gets pin (fastpath) | | | |
| | 1 db block changes | | | |
| | 1 enqueue releases | | | |
| | 1 enqueue requests | | | |
| | 2 execute count | | | |
| | 131072 logical read bytes from cache | | | |
| | 33 non-idle wait count | | | |
| | 2 opened cursors cumulative | | | |
| | 1 opened cursors current | | | |
| | 1 parse count (hard) | | | |
| | 2 parse count (total) | | | |
| | 1 recursive calls | | | |

Picture 1.17 - Count of Used Blocks

--3

SET autotrace ON;

SELECT COUNT(*)

FROM t2 ;

| PLAN_TABLE_OUTPUT | | | | | Statistics | |
|--------------------------------------|-------------------|------|--------|--|--|--|
| SQL_ID 39nv0sqd490k3, child number 0 | | | | | ----- | |
| SELECT COUNT(*) FROM t2 | | | | | 1 CCursor + sql area evicted | |
| Plan hash value: 3321871023 | | | | | 1 CPU used by this session | |
| | | | | | 1 CPU used when call started | |
| | | | | | 1 DB time | |
| | | | | | 33 Requests to/from client | |
| | | | | | 33 SQL*Net roundtrips to/from client | |
| | | | | | 2 buffer is not pinned count | |
| | | | | | 518 bytes received via SQL*Net from client | |
| | | | | | 60749 bytes sent via SQL*Net to client | |
| | | | | | 3 calls to get snapshot scn: kcmgss | |
| | | | | | 3 calls to kcmgcs | |
| | | | | | 1 consistent gets | |
| | | | | | 1 consistent gets from cache | |
| | | | | | 1 consistent gets pin | |
| | | | | | 1 consistent gets pin (fastpath) | |
| | | | | | 1 enqueue releases | |
| | | | | | 1 enqueue requests | |
| | | | | | 2 execute count | |
| | | | | | 32768 logical read bytes from cache | |
| | | | | | 33 non-idle wait count | |
| | | | | | 2 opened cursors cumulative | |
| | | | | | 1 opened cursors current | |
| | | | | | 1 parse count (hard) | |
| | | | | | 2 parse count (total) | |
| | | | | | 1 parse time cpu | |
| | | | | | 1 recursive calls | |
| | | | | | 1 session cursor cache hits | |
| PLAN_TABLE_OUTPUT | | | | | ----- | |
| Id | Operation | Name | E-Rows | | | |
| 0 | SELECT STATEMENT | | | | | |
| 1 | SORT AGGREGATE | | 1 | | | |
| PLAN_TABLE_OUTPUT | | | | | ----- | |
| 2 | TABLE ACCESS FULL | T2 | 99999 | | | |

Picture 1.18 – Quantity of Consistent Gets and Count of Rows

Task Result:

Summary table with all result and text description of analyses these results.

| No | Count of Blocks | Count of Used Blocks | Count of Rows | Consistent gets | Description |
|----|-----------------|----------------------|---------------|-----------------|--|
| 1 | 416 | 376 | 99999 | 380 | Table was filled into 416 blocks, 376 of which were used (40 unused). 380 consistent get operations. Maximal number of rows - 99999 |
| 2 | 416 | 0 | 0 | 380 | We used DELETE operation but still have an opportunity to roll back deleted data. DELETE just saves allocated space and stats, so the number of consistent gets did not change after as you can see below 380 -> 380 |
| 3 | 416 | 1 | 1 | 380 | Even if we INSERT 1 row (or 2,3,4 ... 99999) |
| 4 | 6 | 0 | 0 | 1 | We used TRUNCATE operator which truncate all the allocated data and stats (w/o an opportunity to roll back data), but oracle still leave 6 blocks to store an information about table |

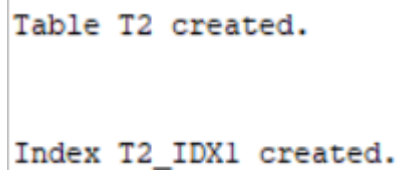
Task 2: Index Clustering factor parameter

Expected:

- Screenshot of the step 5;
- Description of the parameter clustering factor;
- Explanation: why for indexes *t1_idx1* and *t2_idx1* we have different values ;
- Which Index has best selective performance in execution Select clause filtered by IN (, list of values,);

Step 1:

```
CREATE TABLE t2 AS
SELECT TRUNC( rownum / 100 ) id, rpad( rownum,100 ) t_pad
FROM dual
CONNECT BY rownum < 100000;
CREATE INDEX t2_idx1 ON t2
( id );
```



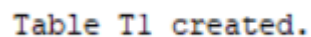
```
Table T2 created.

Index T2_IDX1 created.
```

Picture 2.1 – Table and Index creating

Step 2:

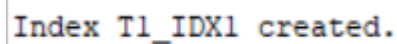
```
# CREATE TABLE t1 AS
SELECT MOD( rownum, 100 ) id, rpad( rownum,100 ) t_pad
FROM dual
CONNECT BY rownum < 100000;
```



```
Table T1 created.
```

Step 3:

```
# CREATE INDEX t1_idx1 ON t1
( id );
```



```
Index T1_IDX1 created.
```

Step 4:

```
# EXEC dbms_stats.gather_table_stats(USER,'t1',method_opt=>'FOR ALL COLUMNS SIZE
1',CASCADE=>TRUE);

# EXEC dbms_stats.gather_table_stats( USER,'t2',method_opt=>'FOR ALL COLUMNS SIZE
1',CASCADE=>TRUE);
```

```
PL/SQL procedure successfully completed.
```

```
PL/SQL procedure successfully completed.
```

Picture 2.2 – Procedure execution

Step 5:

```
# SELECT t.table_name||'.'||i.index_name idx_name,
       i.clustering_factor,
       t.blocks,
       t.num_rows
FROM user_indexes i, user_tables t
WHERE i.table_name = t.table_name
AND t.table_name IN ('T1','T2');
```

| IDX_NAME | CLUSTERING_FACTOR | BLOCKS | NUM_ROWS |
|------------|-------------------|--------|----------|
| T1.T1_IDX1 | 37200 | 381 | 99999 |
| T2.T2_IDX1 | 378 | 386 | 99999 |

Task Results:

- The clustering factor indicates the amount of order of the rows in the table based on the values of the index. At the screenshot above we can see that clustering factor **approaches the number of blocks** in the table 378 ~ 386. That means that the rows are ordered. If it approaches the number of rows in the table, the rows are randomly ordered. In such a case (clustering factor near the number of rows), it is unlikely that index entries in the same leaf block will point to rows in the same data blocks. If all of the index entries in a given leaf block point to different blocks in the table – the table is **not well ordered** with respect to this index. If we want index being **strong** clustered – we should consider using Index Organized Tables (**IOT**). They force the rows into a specific **physical location** based on their index entry. Even though the correct set up of clustering anyway increases performance and decreases cost.
- Because we used MOD operand in creation of table t1 instead of TRUNC in t2.
- Clustering factor increments if rowid points to different table blocks.

Task 3: Index Unique Scan

Step 1:

```
# CREATE UNIQUE INDEX idx_t1 ON t1(t_pad);
```

INDEX UDX_T1 created.

Step 2

```
# SELECT t1.* FROM t1 where t1.t_pad = '1';
```

PLAN_TABLE_OUTPUT

SQL_ID 0bqrmw3lhxhg, child number 0

SELECT t1.* FROM t1 where t1.t_pad = '1'

Plan hash value: 1698228356

| Id | Operation | Name | E-Rows |

| 0 | SELECT STATEMENT | | |
| 1 | TABLE ACCESS BY INDEX ROWID | T1 | 1 |

no rows selected

Statistics

1 CPU used by this session
1 CPU used when call started
1 DB time
33 Requests to/from client
33 SQL*Net roundtrips to/from client
1 buffer is not pinned count
533 bytes received via SQL*Net from client
60509 bytes sent via SQL*Net to client
2 calls to get snapshot scn: kcmgss
2 calls to kcmgss
2 consistent gets
2 consistent gets examination
2 consistent gets examination (fastpath)
2 consistent gets from cache
2 execute count
1 index fetch by kev

Picture 3.1 - Query Result

Note. Select in 3 task is empty because '1' is a char, but a column type is number.

Task 4: Index Range Scan

Step 1:

```
# SELECT t2.* FROM t2 where t2.id = '1';
```

| ID | T_PAD |
|----|-------|
| 1 | 100 |
| 1 | 101 |
| 1 | 102 |
| 1 | 103 |
| 1 | 104 |
| 1 | 105 |
| 1 | 106 |
| 1 | 107 |
| 1 | 108 |
| 1 | 109 |
| 1 | 110 |

Picture 4.1 - Query Result

Task 5: Index Skip Scan

I found scott schema in the internet. Separate file named 'scottschema' will be attached.

Step 1:

```
CREATE TABLE employees AS  
SELECT *  
FROM scott.emp;
```

```
Table EMPLOYEES created.
```

Step 2:

```
CREATE INDEX idx_emp01 ON employees  
( empno, ename, job );
```

```
Index IDX_EMP01 created.
```

Step 3:

```
# SELECT /*+INDEX_SS(emp idx_emp01)*/ emp.* FROM employees emp where ename = 'SCOTT';
```

```
# SELECT /*+FULL*/ emp.* FROM employees emp WHERE ename = 'SCOTT'
```

```
Plan hash value: 1748371071
```

| Id | Operation | Name | Starts | E-Rows | A-Rows | A-Time | Buffers |
|-------------------|-------------------------------------|-----------|--------|--------|--------|-------------|---------|
| 0 | SELECT STATEMENT | | 1 | | 0 | 00:00:00.01 | 1 |
| PLAN_TABLE_OUTPUT | | | | | | | |
| 1 | TABLE ACCESS BY INDEX ROWID BATCHED | EMPLOYEES | 1 | 1 | 0 | 00:00:00.01 | 1 |
| * 2 | INDEX SKIP SCAN | IDX_EMP01 | 1 | 1 | 0 | 00:00:00.01 | 1 |

```
Plan hash value: 1748371071
```

| Id | Operation | Name | Starts | E-Rows | A-Rows | A-Time | Buffers |
|-------------------|-------------------------------------|-----------|--------|--------|--------|-------------|---------|
| 0 | SELECT STATEMENT | | 1 | | 0 | 00:00:00.01 | 1 |
| 1 | TABLE ACCESS BY INDEX ROWID BATCHED | EMPLOYEES | 1 | 1 | 0 | 00:00:00.01 | 1 |
| PLAN_TABLE_OUTPUT | | | | | | | |
| * 2 | INDEX SKIP SCAN | IDX_EMP01 | 1 | 1 | 0 | 00:00:00.01 | 1 |

Note. Oracle scans the index sub-trees for each of the possible values from leading column.

| | Count of Blocks | Count of Used Blocks | Count of Rows | Consistent gets | Description |
|---|-----------------|----------------------|---------------|-----------------|---|
| 1 | 6 | 1 | 1 | 2 | Skip scan because optimizer decided it is more suitable for this SELECT |
| | 6 | 1 | 1 | 2 | /*+INDEX_SS(emp idx_emp01)*/ (cause of hint) |
| | 6 | 1 | 1 | 2 | /*+FULL*/ (cause of hint) |

SELECT BLOCKS FROM user_segments WHERE segment_name = 'EMPLOYEES';

| | BLOCKS |
|---|--------|
| 1 | 6 |

SELECT COUNT(DISTINCT (dbms_rowid.rowid_block_number(rowid))) block_ct FROM EMPLOYEES ;

| | | | | | |
|--|--|--|--|-------|-------------------------------------|
| | | | | | 5 calls to get snapshot scn: kcmgss |
| | | | | | 2 calls to kcmgcs |
| | | | | | 2 consistent gets |
| | | | | | 2 consistent gets from cache |
| | | | | | 2 consistent gets pin |
| | | | | | 2 consistent gets pin (fastpath) |
| | | | | | 1 enqueue releases |
| | | | | | 1 enqueue requests |
| | | | | | 2 execute count |
| | | | | 65536 | logical read bytes from cache |
| | | | | 2 | no work - consistent read gets |
| | | | | 43 | non-idle wait count |
| | | | | 2 | opened cursors cumulative |
| | | | | 1 | opened cursors current |
| | | | | 1 | parse count (hard) |
| | | | | 2 | parse count (total) |
| | | | | 1 | parse time elapsed |

| | BLOCK_CT | | BLOCK_CT |
|---|----------|---|----------|
| 1 | 1 | 1 | 1 |

Review

Laboratory Work Summary:

In this laboratory work we have learned and practiced about how Oracle uses different techniques to access the data needed for the query answer. Access Methods can be divided into two categories: data accessed via a table scan or index access. There are some description below about what methods have we used to access the data at this lab:

Full Scan:

A full table scan reads all rows from a table, and then filters out those rows that do not meet the selection criteria.

In a full table scan, the database sequentially reads every formatted block under the high water mark. The database reads each block only once.

Clustering factor:

The clustering factor is a number that represents the degree to which data is randomly distributed in a table as compared to the indexed column. In simple terms, it is the number of "block switches" while reading a table using an index.

Index Unique Scan:

This scan returns, at most, a single rowid. Oracle performs a unique scan if a statement contains a UNIQUE or a PRIMARY KEY constraint that guarantees that only a single row is accessed.

Index Range Scan:

An index range scan is a common operation for accessing selective data. It can be bounded (bounded on both sides) or unbounded (on one or both sides). Data is returned in the ascending order of index columns. Multiple rows with identical values are sorted in ascending order by rowid.

Index Skip Scan:

Index skip scans improve index scans by non-prefix columns. Often, scanning index blocks is faster than scanning table data blocks. Skip scanning lets a composite index be split logically into smaller sub-indexes. In skip scanning, the initial column of the composite index is not specified in the query. In other words, it is skipped. The number of logical sub-indexes is determined by the number of distinct values in the initial column. Skip scanning is advantageous if there are few different values in the leading column of the composite index and many distinct values in the non-leading key of the index.