

# Report

Laboratory Work 4

Dmitry Ladutsko

July 19, 2022

## Task 1: Auto Trace configuration training

Nº	Auto Trace Configuration Options	Expected Results	Description
	set autotrace off	No AUTOTRACE report is generated.	This is the default.
	set autotrace on	The AUTOTRACE report	Includes both the optimizer execution path and the SQL statement execution statistics.
	set autotrace traceonly	Like SET AUTOTRACE ON	Suppresses the printing of the user's query output, if any. This option is useful when you are tuning a large query, but do not want to see the query report.
	set autotrace on explain	The AUTOTRACE report	Shows only the optimizer execution path.
	set autotrace on statistics	The AUTOTRACE report	Shows only the SQL statement execution statistics.
	set autotrace on explain statistics	Shows only the SQL statement execution statistics.	= set autotrace on
	set autotrace traceonly explain	Execution plan w/o query result	These options should be used when a large result set is expected.
	set autotrace traceonly statistics	Statistics w/o query result	
	set autotrace traceonly explain statistics	Execution plan + statistics w/o query result	
	set autotrace off explain	Disables autotrace utility	
	set autotrace off statistics		
	set autotrace off explain statistics		

## Task 2: Nested Loop Joins

### Example:

```
SELECT /*+ gather_plan_statistics */ *
```

```
FROM scott.emp e, scott.dept d
```

```
WHERE e.deptno = d.deptno
```

```
AND d.deptno = 10
```

Id	Operation	Name	Starts	E-Rows	A-Rows	A-Time	Buffers
0	SELECT STATEMENT		1		3	100:00:00.01	6

PLAN\_TABLE\_OUTPUT

1	NESTED LOOPS		1	5	3	100:00:00.01	6
2	TABLE ACCESS BY INDEX ROWID	DEPT	1	1	1	100:00:00.01	2
* 3	INDEX UNIQUE SCAN	PK_DEPT	1	1	1	100:00:00.01	1
* 4	TABLE ACCESS FULL	EMP	1	5	3	100:00:00.01	4

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				5
NESTED LOOPS				5
TABLE ACCESS	SCOTT.DEPT	BY INDEX ROWID		1
INDEX	SCOTT.PK_DEPT	UNIQUE SCAN		1
Access Predicates				
D.DEPTNO=10				
TABLE ACCESS	SCOTT.EMP	FULL		5
Filter Predicates				
E.DEPTNO=10				

Picture 2.1 - Query Result

*Note.* I used `/*+ gather_plan_statistics */` hint because for some reason explain plan statement did not give me full needed stats (only id, operation and name columns). Nested loops join compare every row of table B with current row in table A. If it is a row, which satisfies the specified condition, it is selected, otherwise goes next and compare second row from table A with all rows from table B. This method is quite useful with small table, but not with big.

## Task 3: Sort-Merge Joins

```
set autotrace on explain;
```

```
SELECT /*+ use_merge(e d) gather_plan_statistics */ *
```

```
FROM scott.emp e, scott.dept d
```

```
WHERE e.deptno = d.deptno
```

Id	Operation	Name	Starts	E-Rows	A-Rows	A-Time	Buffers	OMem	lMem	Used-Mem
0	SELECT STATEMENT		1		14	00:00:00.01	6			
PLAN_TABLE_OUTPUT										
1	MERGE JOIN		1	14	14	00:00:00.01	6			
2	TABLE ACCESS BY INDEX ROWID	DEPT	1	4	4	00:00:00.01	2			
3	INDEX FULL SCAN	PK_DEPT	1	4	4	00:00:00.01	1			
* 4	SORT JOIN		4	14	14	00:00:00.01	4	2048	2048	2048 (0)
5	TABLE ACCESS FULL	EMP	1	14	14	00:00:00.01	4			

Picture 3.1 - Query Result

**Note.** Using this type of join we need to be sure that all used tables are sorted by columns, mentioned in query condition. Oracle scans each row once. This gives us nice benefits, but we should keep in mind that sorting loads the system a lot if we use big tables.

## Task 4: Hash Joins

set autotrace on explain;

SELECT /\*+ gather\_plan\_statistics USE\_HASH(e d) \*/ \*

FROM scott.emp e, scott.dept d

WHERE e.deptno = d.deptno

Id	Operation	Name	Starts	E-Rows	A-Rows	A-Time	Buffers	OMem	lMem	Used-Mem
0	SELECT STATEMENT		1		14	00:00:00.01	36			
PLAN_TABLE_OUTPUT										
* 1	HASH JOIN		1	14	14	00:00:00.01	36	801K	801K	692K (0)
2	TABLE ACCESS FULL	DEPT	1	4	4	00:00:00.01	4			
3	TABLE ACCESS FULL	EMP	1	14	14	00:00:00.01	4			

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				14 8
HASH JOIN				14 8
Access Predicates				
E.DEPTNO=D.DEPTNO				
TABLE ACCESS	SCOTT.DEPT	FULL	4	4
TABLE ACCESS	SCOTT.EMP	FULL	14	4

Picture 4.1 - Query Result

**Note.** Hash join is used when projections of the joined tables are not sorted on the join columns. In this case, the optimizer builds an in-memory hash table on the inner table's join column. The optimizer then scans the outer table for matches to the hash table, and joins data from the two tables accordingly.

## Task 5: Cartesian Joins

```
set autotrace on explain;
```

```
SELECT /*+ gather_plan_statistics */ *
```

```
FROM scott.emp, scott.dept
```

Id	Operation	Name	Starts	E-Rows	A-Rows	A-Time	Buffers	OMem	lMem	Used-Mem
0	SELECT STATEMENT		1		56	00:00:00.01	9			
1	MERGE JOIN CARTESIAN		1	56	56	00:00:00.01	9			
PLAN_TABLE_OUTPUT										
2	TABLE ACCESS FULL	DEPT	1	4	4	00:00:00.01	5			
3	BUFFER SORT		4	14	56	00:00:00.01	4	2048	2048	2048 (0)
4	TABLE ACCESS FULL	EMP	1	14	14	00:00:00.01	4			

  

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				56 12
MERGE JOIN		CARTESIAN		56 12
TABLE ACCESS	SCOTT.DEPT	FULL		4 4
BUFFER		SORT		14 8
TABLE ACCESS	SCOTT.EMP	FULL		14 2

Picture 5.1 - Query Result

*Note.* Cartesian joins are used when all rows in all tables listed in a query: each row in the first table is paired with all the rows in the second table. This happens when there is no relationship defined between the two tables.

## Task 6: Left/Right Outer Joins

### 1. Left outer JOIN

```
set autotrace on explain;
```

```
select /*+ gather_plan_statistics */ e.ename, e.deptno, e.job, d.dname
```

```
from scott.emp e, scott.dept d
```

```
where e.deptno = d.deptno(+);
```

Id	Operation	Name	Starts	E-Rows	A-Rows	A-Time	Buffers	OMem	lMem	Used-Mem
0	SELECT STATEMENT		1		14	00:00:00.01	8			
PLAN_TABLE_OUTPUT										
* 1	HASH JOIN OUTER		1	14	14	00:00:00.01	8	830K	830K	692K (0)
2	TABLE ACCESS FULL	EMP	1	14	14	00:00:00.01	4			
3	TABLE ACCESS FULL	DEPT	1	4	4	00:00:00.01	4			

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				14
HASH JOIN		OUTER		14
Access Predicates E.DEPTNO=D.DEPTNO(+)				8
NESTED LOOPS		OUTER		14
STATISTICS COLLECTOR				8
TABLE ACCESS	SCOTT.EMP	FULL		14
TABLE ACCESS	SCOTT.DEPT	BY INDEX ROWID	1	4
INDEX	SCOTT.PK_DEPT	UNIQUE SCAN		4
Access Predicates E.DEPTNO=D.DEPTNO(+)				
TABLE ACCESS	SCOTT.DEPT	FULL	4	4

Picture 6.1 - Query Result

**Note.** The specific of this type of join is that result includes unmatched rows from only the table that is specified before the LEFT OUTER JOIN clause. If we are joining two tables and want the result set to include unmatched rows from only one table, we use a LEFT OUTER JOIN clause (or a RIGHT OUTER JOIN) clause.

## 2. Right outer JOIN

set autotrace on explain;

select /\*+ gather\_plan\_statistics \*/ e.ename, e.deptno, e.job, d.dname

from scott.emp e, scott.dept d

where e.deptno(+) = d.deptno;

Id	Operation	Name	Starts	E-Rows	A-Rows	A-Time	Buffers	OMem	lMem	Used-Mem
0	SELECT STATEMENT		1		15	00:00:00.01	6			
PLAN_TABLE_OUTPUT										
1	MERGE JOIN OUTER		1	15	15	00:00:00.01	6			
2	TABLE ACCESS BY INDEX ROWID	DEPT	1	4	4	00:00:00.01	2			
3	INDEX FULL SCAN	PK_DEPT	1	4	4	00:00:00.01	1			
* 4	SORT JOIN		4	14	14	00:00:00.01	4	2048	2048	2048 (0)
5	TABLE ACCESS FULL	EMP	1	14	14	00:00:00.01	4			

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				15
MERGE JOIN		OUTER		15
TABLE ACCESS	SCOTT.DEPT	BY INDEX ROWID	4	2
INDEX	SCOTT.PK_DEPT	FULL SCAN	4	1
JOIN				14
Access Predicates E.DEPTNO(+) = D.DEPTNO				
Filter Predicates E.DEPTNO(+) = D.DEPTNO				
TABLE ACCESS	SCOTT.EMP	FULL	14	4

Picture 6.2 - Result

## Task 7: Full Outer Join

set autotrace on explain;

select /\*+ gather\_plan\_statistics \*/ e.ename, e.deptno, e.job, d.dname

from scott.emp e

full outer join

scott.dept d

on(e.deptno = d.deptno);

Id	Operation	Name	Starts	E-Rows	A-Rows	A-Time	Buffers	OMem	lMem	Used-Mem
PLAN_TABLE_OUTPUT										
0	SELECT STATEMENT		1		15	00:00:00.01	8			
1	VIEW	VW_FOJ_0	1	15	15	00:00:00.01	8			
* 2	HASH JOIN FULL OUTER		1	15	15	00:00:00.01	8	971K	971K	969K (0)
3	TABLE ACCESS FULL	DEPT	1	4	4	00:00:00.01	4			
4	TABLE ACCESS FULL	EMP	1	14	14	00:00:00.01	4			

  

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				15
VIEW	SYS.VW_FOJ_0			15
HASH JOIN		FULL OUTER		15
Access Predicates				
E.DEPTNO=D.DEPTNO				
TABLE ACCESS	SCOTT.DEPT	FULL		4
TABLE ACCESS	SCOTT.EMP	FULL		14

Picture 7.1 - Result

**Note.** We use Full outer join if we want the result set to include unmatched rows from both tables.

## Task 8: Semi Joins

### 1. Using /\*semijoin\*/:

set autotrace on explain;

select /\*+ gather\_plan\_statistics semijoin \*/ dname

from scott.dept d

where deptno in (select deptno from scott.emp e);

Id	Operation	Name	Starts	E-Rows	A-Rows	A-Time	Buffers	OMem	lMem	Used-Mem
PLAN_TABLE_OUTPUT										
0	SELECT STATEMENT		1		3	00:00:00.01	6			
1	MERGE JOIN SEMI		1	3	3	00:00:00.01	6			
2	TABLE ACCESS BY INDEX ROWID	DEPT	1	4	4	00:00:00.01	2			
3	INDEX FULL SCAN	PK_DEPT	1	4	4	00:00:00.01	1			
* 4	SORT UNIQUE		4	14	3	00:00:00.01	4	2048	2048	2048 (0)
5	TABLE ACCESS FULL	EMP	1	14	14	00:00:00.01	4			

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				7
MERGE JOIN		SEMI		7
TABLE ACCESS	SCOTT.DEPT	BY INDEX ROWID		2
INDEX	SCOTT.PK_DEPT	FULL SCAN		1
SORT		UNIQUE	14	5
Access Predicates				
DEPTNO=DEPTNO				
Filter Predicates				
DEPTNO=DEPTNO				
TABLE ACCESS	SCOTT.EMP	FULL	14	4

Picture 8.1 -Query Result

**Note.** A semi-join returns one copy of each row in first table for which at least one match is found. Semi-joins are also written using the EXISTS construct.

## 2. Using /\*no semijoin\*/:

set autotrace on explain;

select /\*+ gather\_plan\_statistics no\_semijoin \*/ dname

from scott.dept d

where deptno in (select deptno from scott.emp e);

Id	Operation	Name	Starts	E-Rows	A-Rows	A-Time	Buffers	OMem	lMem	Used-Mem
0	SELECT STATEMENT		1		3	00:00:00.01	6			
PLAN_TABLE_OUTPUT										
1	MERGE JOIN SEMI		1	3	3	00:00:00.01	6			
2	TABLE ACCESS BY INDEX ROWID	DEPT	1	4	4	00:00:00.01	2			
3	INDEX FULL SCAN	PK_DEPT	1	4	4	00:00:00.01	1			
4	SORT UNIQUE		4	14	3	00:00:00.01	4	2048	2048	2048 (0)
5	TABLE ACCESS FULL	EMP	1	14	14	00:00:00.01	4			

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				7
MERGE JOIN		SEMI		7
TABLE ACCESS	SCOTT.DEPT	BY INDEX ROWID		2
INDEX	SCOTT.PK_DEPT	FULL SCAN		1
SORT		UNIQUE	14	5
Access Predicates				
DEPTNO=DEPTNO				
Filter Predicates				
DEPTNO=DEPTNO				
TABLE ACCESS	SCOTT.EMP	FULL	14	4

Picture 8.2 -Query Result

**Note.** As a consequence, NO vital difference

## 3. Using /\*using exists\*/

set autotrace on explain;

select /\*+ gather\_plan\_statistics \*/ dname

from scott.dept d

where exists (select null from scott.emp e

where e.deptno = d.deptno);



Id	Operation	Name	Starts	E-Rows	A-Rows	A-Time	Buffers	OMem	lMem	Used-Mem
----	-----------	------	--------	--------	--------	--------	---------	------	------	----------

  

PLAN_TABLE_OUTPUT										
0	SELECT STATEMENT		1		3	00:00:00.01	6			
1	MERGE JOIN SEMI		1	3	3	00:00:00.01	6			
2	TABLE ACCESS BY INDEX ROWID	DEPT	1	4	4	00:00:00.01	2			
3	INDEX FULL SCAN	PK_DEPT	1	4	4	00:00:00.01	1			
* 4	SORT UNIQUE		4	14	3	00:00:00.01	4	2048	2048	2048 (0)
5	TABLE ACCESS FULL	EMP	1	14	14	00:00:00.01	4			

  

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				7
MERGE JOIN		SEMI		7
TABLE ACCESS	SCOTT.DEPT	BY INDEX ROWID		2
INDEX	SCOTT.PK_DEPT	FULL SCAN		1
SORT		UNIQUE		5
Access Predicates				
E.DEPTNO=D.DEPTNO				
Filter Predicates				
E.DEPTNO=D.DEPTNO				
TABLE ACCESS	SCOTT.EMP	FULL		4

Picture 8.3 -Query Result

**Note.** The EXISTS function checks to find a single matching row to return the result in a subquery. Because the IN function retrieves and checks all rows, so it is slower.

#### 4. Select distinct using /\* inner join with distinct \*/ hint

set autotrace on explain;

select /\*+ gather\_plan\_statistics inner join with distinct \*/ distinct dname

from scott.dept d,scott.emp e

where d.deptno = e.deptno ;

Id	Operation	Name	Starts	E-Rows	A-Rows	A-Time	Buffers	OMem	lMem	Used-Mem
----	-----------	------	--------	--------	--------	--------	---------	------	------	----------

  

PLAN_TABLE_OUTPUT										
1	HASH UNIQUE		1	3	3	00:00:00.01	6	714K	714K	
2	MERGE JOIN SEMI		1	3	3	00:00:00.01	6			
3	TABLE ACCESS BY INDEX ROWID	DEPT	1	4	4	00:00:00.01	2			
4	INDEX FULL SCAN	PK_DEPT	1	4	4	00:00:00.01	1			
* 5	SORT UNIQUE		4	14	3	00:00:00.01	4	2048	2048	2048 (0)
6	TABLE ACCESS FULL	EMP	1	14	14	00:00:00.01	4			

  

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				8
HASH		UNIQUE		8
MERGE JOIN		SEMI		7
TABLE ACCESS	SCOTT.DEPT	BY INDEX ROWID		2
INDEX	SCOTT.PK_DEPT	FULL SCAN		1
SORT		UNIQUE		5
Access Predicates				
D.DEPTNO=E.DEPTNO				
Filter Predicates				
D.DEPTNO=E.DEPTNO				
TABLE ACCESS	SCOTT.EMP	FULL		4

Picture 8.4 -Query Result

**Note.** The DISTINCT clause is used in a SELECT statement to filter duplicate rows in the result set. This way we can ensure that returned rows are unique for the column or columns specified in the SELECT clause.

## Task 9: Anti Joins

### 1. Not exists using /\*not exists\*/ hint

set autotrace on explain;

select /\*+ gather\_plan\_statistics not exists ANTIJOIN \*/ dname

from scott.dept d

where not exists (select null from scott.emp e

where e.deptno = d.deptno);

Id	Operation	Name	Starts	E-Rows	A-Rows	A-Time	Buffers	OMem	lMem	Used-Mem
PLAN_TABLE_OUTPUT										
0	SELECT STATEMENT		1		1	00:00:00.01	6			
1	MERGE JOIN ANTI		1	1	1	00:00:00.01	6			
2	TABLE ACCESS BY INDEX ROWID	DEPT	1	4	4	00:00:00.01	2			
3	INDEX FULL SCAN	PK_DEPT	1	4	4	00:00:00.01	1			
* 4	SORT UNIQUE		4	14	3	00:00:00.01	4	2048	2048	2048 (0)
5	TABLE ACCESS FULL	EMP	1	14	14	00:00:00.01	4			

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				7
MERGE JOIN		ANTI		7
TABLE ACCESS	SCOTT.DEPT	BY INDEX ROWID	4	2
INDEX	SCOTT.PK_DEPT	FULL SCAN	4	1
SORT		UNIQUE		5
Access Predicates				
E.DEPTNO=D.DEPTNO				
Filter Predicates				
E.DEPTNO=D.DEPTNO				
TABLE ACCESS	SCOTT.EMP	FULL	14	4

**Note.** The NOT EXISTS operator returns true if the subquery returns no row. Otherwise, it returns false. We should keep in mind that NOT EXISTS operator returns false if the subquery returns any rows with a NULL value.

### 2. Not in using /\*not in\*/ hint

set autotrace on explain;

select /\*+ gather\_plan\_statistics not in \*/ dname

from scott.dept d

where deptno not in

(select deptno from scott.emp e)

Id	Operation	Name	Starts	E-Rows	A-Rows	A-Time	Buffers	OMem	lMem	Used-Mem
PLAN_TABLE_OUTPUT										
0	SELECT STATEMENT		1			00:00:00.01	7			
1	MERGE JOIN ANTI NA		1	1	1	00:00:00.01	7			
2	SORT JOIN		1	4	4	00:00:00.01	0	2048	2048	2048 (0)
3	TABLE ACCESS BY INDEX ROWID BATCHED	DEPT	1	4	4	00:00:00.01	2			
4	INDEX FULL SCAN	PK_DEPT	1	4	4	00:00:00.01	1			
* 5	SORT UNIQUE		5	14	3	00:00:00.01	0	2048	2048	2048 (0)
6	TABLE ACCESS FULL	EMP	1	14	14	00:00:00.01	5			

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				7
MERGE JOIN		ANTI NA		7
SORT		JOIN		2
TABLE ACCESS	SCOTT.DEPT	BY INDEX ROWID BATCHED		2
INDEX	SCOTT.PK_DEPT	FULL SCAN		1
SORT		UNIQUE		5
Access Predicates				
DEPTNO=DEPTNO				
Filter Predicates				
DEPTNO=DEPTNO				
TABLE ACCESS	SCOTT.EMP	FULL		4

*Note.* Using NOT IN operator can be risky because it does not return null values

NOT IN can be just as efficient as NOT EXISTS even if an anti-join can be used (if the subquery is known to not return nulls)

### 3. Minus using /\*minus\*/ hint

set autotrace on explain;

select /\*+ gather\_plan\_statistics minus\*/ dname

from scott.dept

where deptno in

(select deptno from scott.dept minus

select deptno from scott.dept);

Id	Operation	Name	Starts	E-Rows	A-Rows	A-Time	Buffers	OMem	lMem	Used-Mem
PLAN_TABLE_OUTPUT										
0	SELECT STATEMENT		1		0	00:00:00.01	4			
1	MERGE JOIN		1	4	0	00:00:00.01	4			
2	TABLE ACCESS BY INDEX ROWID	DEPT	1	4	1	00:00:00.01	2			
3	INDEX FULL SCAN	PK_DEPT	1	4	1	00:00:00.01	1			
4	SORT JOIN		1	4	0	00:00:00.01	2	1024	1024	
5	VIEW	VW_NSO_1	1	4	0	00:00:00.01	2			
6	MINUS		1		0	00:00:00.01	2			
7	SORT UNIQUE		1	4	4	00:00:00.01	1	2048	2048	2048 (0)
8	INDEX FULL SCAN	PK_DEPT	1	4	4	00:00:00.01	1			
9	SORT UNIQUE		1	4	4	00:00:00.01	1	2048	2048	2048 (0)
10	INDEX FULL SCAN	PK_DEPT	1	4	4	00:00:00.01	1			

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				7
MERGE JOIN				7
TABLE ACCESS	SCOTT.DEPT	BY INDEX ROWID		2
INDEX	SCOTT.PK_DEPT	FULL SCAN		1
SORT		JOIN		5
Access Predicates				
DEPTNO=DEPTNO				
Filter Predicates				
DEPTNO=DEPTNO				
VIEW	SYS.VW_NSO_1			4
MINUS				
SORT		UNIQUE		4
INDEX	SCOTT.PK_DEPT	FULL SCAN		1
SORT		UNIQUE		4
INDEX	SCOTT.PK_DEPT	FULL SCAN		1

*Note.* The Oracle MINUS operator is used to return all rows in the first SELECT statement that are not returned by the second SELECT statement. Each SELECT statement will define a dataset. The MINUS operator retrieve all records from the first dataset and then remove from the results all records from the second dataset.

#### 4. Left outer anti join

set autotrace on explain;

select /\*+ gather\_plan\_statistics left outer \*/ dname

from scott.dept d, scott.emp e

where d.deptno = e.empno(+)

and e.deptno is null;

Id	Operation	Name	Starts	E-Rows	A-Rows	A-Time	Buffers
0	SELECT STATEMENT		1		4	100:00:00.01	6

  

Id	Operation	Name	Starts	E-Rows	A-Rows	A-Time	Buffers
1	FILTER		1		4	100:00:00.01	6
2	NESTED LOOPS OUTER		1	4	4	100:00:00.01	6
3	TABLE ACCESS FULL	DEPT	1	4	4	100:00:00.01	4
4	TABLE ACCESS BY INDEX ROWID	EMP	4	1	0	100:00:00.01	2
5	INDEX UNIQUE SCAN	PK_EMP	4	1	0	100:00:00.01	2

  

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				4
FILTER				
Filter Predicates				
E.DEPTNO IS NULL				
NESTED LOOPS				
TABLE ACCESS	SCOTT.DEPT	OUTER		4
TABLE ACCESS	SCOTT.EMP	FULL		4
INDEX	SCOTT.PK_EMP	BY INDEX ROWID		1
Access Predicates				
D.DEPTNO=E.EMPNO(+)		UNIQUE SCAN		1

Picture 9.1 -Query Result

### Task 10: Summary table

Join Access “A”	Join Access “B”	Nested Loop	Hash Join	Sort-Merge Join	Anti-Join	Semi-Join
Small Table	Small Table	Good	Ineffective	Ineffective	Contextual	
Small Table	Small Table(Indexed)			Good		
Small Table(Indexed)	Small Table(Indexed)			Good		
Big table	Big table	Awful	Good	Ineffective		
Big table(Indexed)	Big table			Good		
Big table(Indexed)	Big table(Indexed)			Good		
Small	Big		Ineffective			

## Laboratory work summary:

We used many join methods to see which has better performance in different situations.

At the table above and insights from each task (or sub – task) you can see some results and conclusion.

